

Grouper: Optimal Assignment Workflows for Higher Education

by Mingyuan Zhang, Kevin Lam, and Vik Gopal

Abstract Grouper is an R package for formulating and solving allocation problems that arise in higher education. The package provides three mixed-integer optimization workflows. They cover diversity-based assignment of students to groups and topics, preference-based assignment of self-formed groups to topic slots, and multi-role workload allocation under fairness, preference, and priority constraints. Existing optimization tools in R provide powerful low-level modeling and solver interfaces, but instructors and administrators often need formulations that connect institutional rules to interpretable allocation outputs. Grouper addresses this gap through policy-aware mathematical models, structured input formats, configurable constraints and objective weights, and a stable API pattern across assignment settings. Users can adapt group-size limits, topic capacities, preference weights, workload caps, fairness terms, and priority rules without rewriting the underlying model code. This paper presents the models, package design, implementation, and a detailed workflow demonstration, illustrating how Grouper can support reproducible and transparent allocation decisions in higher-education institutions.

1 Introduction

Allocation problems are a recurring feature of higher education. Instructors and administrators routinely need to assign students to groups, topics, supervisors, teaching duties, or other limited resources while balancing preferences, capacity, fairness, and institutional rules. These decisions often involve several competing objectives. For example, a course may need diverse project groups, a program may need to match self-formed groups to topic slots, and a department may need to distribute teaching or support duties fairly across people.

In many cases, the allocation policy is clear in principle but difficult to operationalize. Manual allocation can be time-consuming, difficult to reproduce, and hard to audit when multiple constraints and policy goals interact.

The grouper package provides a high-level R workflow for optimization-based allocation in higher education. It currently supports three policy-aware workflows. These address diversity-based assignment of students to groups and topics, preference-based assignment of self-formed groups to topic slots, and multi-role workload allocation.

The contribution of grouper is both mathematical and computational. Mathematically, it encodes recurring allocation priorities such as diversity, preference satisfaction, capacity control, workload fairness, and priority-aware assignment as explicit objectives and constraints. Computationally, it exposes these models through a common R workflow in which users prepare structured data, specify policy parameters, solve the model through existing R optimization infrastructure, and inspect allocation results. This design allows users to adapt allocation rules to different courses, cohorts, departments, or workload settings without rewriting the underlying model code, making the allocation process more reproducible and transparent than ad hoc manual assignment.

This paper introduces the mathematical models, implementation, and main workflows of grouper. The next section reviews related work, including general-purpose R optimization tools and existing higher-education allocation approaches. We then present the three allocation models, describe how the package implements them, and use the multi-role workload workflow as the detailed reproducible demonstration on departmental planning data.

Related work

Optimization-based allocation has a long history in higher-education planning. Student-project allocation, group formation, and workload distribution can often be expressed as mathematical programming problems in which assignment decisions are represented by binary or integer variables and institutional rules are encoded as constraints. Prior work has used integer programming approaches to assign students to projects or topics while accounting for preferences, capacities, and other institutional constraints (Anwar and Bahaj, 2003; Meyer, 2009; Ramotsisi et al., 2022).

In R, users can access a rich ecosystem of general-purpose optimization tools. Packages such as `ompr` provide algebraic modeling interfaces for mixed-integer programming (Schumacher, 2025a), while `ROI` provides a common interface to multiple solver plugins (Hornik et al., 2024; Schumacher,

2025b).

Solver backends such as Gurobi, GLPK, and HiGHS can then be used to solve the resulting optimization problems (Gurobi Optimization, LLC, 2026; Free Software Foundation, 2026; HiGHS Developers, 2026). These tools are powerful and flexible, but they operate at a relatively low level of abstraction. Users must still define the decision variables, constraints, objective functions, data transformations, and post-processing steps for each allocation problem.

Several tools are closer to higher-education assignment. Student-project allocation has been studied through integer programming, and web-based or standalone implementations exist for specific versions of the problem (Anwar and Bahaj, 2003; Meyer, 2009). On CRAN, `golfr` provides classroom group assignment functionality with a focus on minimizing repeated pairings across multiple rounds (Kim and Nolte, 2025). These tools address important subproblems, but they are generally designed around a single assignment setting or algorithmic goal.

`grouper` differs from these approaches by providing a higher-level package interface for multiple higher-education allocation workflows. The package does not replace lower-level optimization software; instead, it packages recurring models into documented workflows that combine structured input preparation, configurable policy parameters, mixed-integer model construction, solver execution, and interpretable output generation. To our knowledge, no existing R package provides a comparable orchestration layer for diversity-based student grouping, preference-based topic allocation, and multi-role workload allocation within one extensible framework.

2 Mathematical allocation models

Each `grouper` workflow is formulated as an integer or mixed-integer program. The objective describes the allocation goal, while the constraints define which allocations are feasible. The notation below follows the package vignettes so that the mathematical description can be compared directly with the detailed model documentation.

Diversity-based assignment

Purpose and inputs

Consider N students arranged into G self-formed groups. Each group is assigned to one of T topics, and topic t may have R_t repetitions. The model is applicable to tutorial sections, lab teams, interdisciplinary discussion groups, and capstone teams where instructors want heterogeneous final groups.

The user supplies the student data, the self-formed-group column, demographic information or a dissimilarity matrix, and optionally a numeric skill measure. The extractor constructs m_{ig} , which equals one when student i belongs to self-formed group g , the pairwise dissimilarity values d_{ij} , and the optional skill values s_i .

Objective

$$\max \quad w_1 \left(\sum_{i=1}^{N-1} \sum_{j=i+1}^N \sum_{t=1}^T \sum_{r=1}^{R_t} z_{ijtr} d_{ij} \right) + w_2 (s_{\min} - s_{\max}).$$

Here z_{ijtr} indicates whether students i and j are assigned to the same repetition r of topic t . The first term maximizes pairwise dissimilarity within the final groups. The second minimizes the difference between their maximum and minimum total skill. The non-negative weights w_1 and w_2 determine the relative priority of the two terms and are normally chosen to sum to one. If skill is not supplied, the second term is omitted.

Key constraints

Let

$$x_{gtr} = \begin{cases} 1, & \text{if group } g \text{ is assigned to repetition } r \text{ of topic } t, \\ 0, & \text{otherwise.} \end{cases}$$

Each self-formed group must be assigned exactly once:

$$\sum_{t=1}^T \sum_{r=1}^{R_t} x_{gtr} = 1, \quad \forall g. \quad (1)$$

This equality preserves every self-formed group as a unit. It rules out both leaving a group unassigned and assigning the same group to multiple topic-repetition combinations.

The binary variable z_{ijtr} records whether students i and j are both assigned to the topic-repetition combination:

$$\begin{aligned} z_{ijtr} &\leq \sum_{g=1}^G m_{ig} x_{gtr}, \\ z_{ijtr} &\leq \sum_{g=1}^G m_{jg} x_{gtr}, \quad \forall i < j, t, r. \\ z_{ijtr} &\geq \sum_{g=1}^G m_{ig} x_{gtr} + \sum_{g=1}^G m_{jg} x_{gtr} - 1, \end{aligned} \quad (2)$$

The first two inequalities prevent z_{ijtr} from being one unless both students are in the slot. The third forces it to one when both are present. Consequently, d_{ij} contributes to the objective exactly when the student pair appears in the same final group.

Let $a_{tr} = 1$ when repetition r of topic t is active. The following constraints link activation to assignment and bound the number of repetitions:

$$\begin{aligned} a_{tr} &\geq x_{gtr}, \quad \forall g, t, r, \\ a_{tr} &\leq \sum_{g=1}^G x_{gtr}, \quad \forall t, r, \\ \sum_{r=1}^{R_t} a_{tr} &\geq r_{\min}, \quad \forall t, \\ \sum_{r=1}^{R_t} a_{tr} &\leq r_{\max}, \quad \forall t. \end{aligned} \quad (3)$$

The first two conditions make a_{tr} an indicator for whether the slot is used. The final two then count the used repetitions of each topic and keep that count between the instructor-specified limits $r_{\min}$ and $r_{\max}$.

The number of students in each active topic-repetition combination is bounded:

$$\begin{aligned} \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} &\geq a_{tr} n_{tr}^{\min}, \\ \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} &\leq a_{tr} n_{tr}^{\max}, \end{aligned} \quad \forall t, r. \quad (4)$$

Multiplication by a_{tr} switches the bounds off for an inactive slot. For an active slot, the total number of students contributed by its assigned self-formed groups must lie between $n_{tr}^{\min}$ and $n_{tr}^{\max}$.

When skill balancing is used, $s_{\min}$ and $s_{\max}$ bound the total skill assigned to each active topic-repetition combination:

$$\begin{aligned} \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} s_i &\geq s_{\min}, \\ \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} s_i &\leq s_{\max}, \end{aligned} \quad \forall t, r \text{ with } a_{tr} = 1. \quad (5)$$

These bounds make $s_{\min}$ and $s_{\max}$ the lowest and highest total skill among active slots. Maximizing $s_{\min} - s_{\max}$ therefore minimizes the skill spread. Inactive candidate repetitions are excluded from this comparison.

Table 1: Parameter effect guide for the diversity-based assignment model (directional, holding other settings fixed).

parameter	if increased	trade-off
w1	places greater relative emphasis on pairwise diversity	may place less emphasis on skill balance
w2	places greater relative emphasis on reducing skill differences	may accept lower pairwise diversity
nmin	requires more students in each active topic-repetition combination	tightens feasibility
nmax	permits more students in each topic-repetition combination	increases flexibility but may permit less even group sizes
rmin	requires more active repetitions of every topic	tightens feasibility and spreads students across more repetitions
rmax	permits more active repetitions of every topic	increases flexibility but may fragment the cohort
n_topics	requires assignments across more topics when rmin is positive	increases model size and may tighten feasibility
R (R_t)	provides more candidate repetitions	increases model size and only adds flexibility when repetition bounds permit their use

Preference-based assignment

Purpose and inputs

Unlike the diversity-based model, which assigns groups according to student attributes and the desired composition of the final slots, this model is used when groups have expressed preferences over the available destinations. The self-formed groups remain intact, and the main decision is where each group should be placed rather than which student profiles should be brought together.

Consider N students arranged into G self-formed groups. Each group is assigned to a topic t from T base topics. A topic may contain B subgroups, giving BT topic-subgroup combinations, and each combination may be repeated R_t times. This setting arises in the allocation of project topics, case studies, supervisors, lab rotations, presentation themes, and tutorial or project streams.

The user supplies the student data, self-formed-group column, and preference matrix. The extractor obtains m_{ig} , group sizes n_g , and preference scores p_{tg} , where higher values indicate stronger preferences.

Objective

$$\max \sum_{g=1}^G \sum_{t=1}^{BT} \sum_{r=1}^{R_t} x_{gtr} n_g p_{tg}.$$

Here $x_{gtr} = 1$ when group g is assigned to repetition r of topic-subgroup combination t . Multiplication by n_g means that an assignment affecting a larger group contributes more to total student-level preference satisfaction.

Key constraints

Every self-formed group must be assigned exactly once:

$$\sum_{t=1}^{BT} \sum_{r=1}^{R_t} x_{gtr} = 1, \quad \forall g. \quad (6)$$

As in the diversity model, this equality keeps each self-formed group intact and assigns it to one, and only one, topic-subgroup repetition.

Let $a_{tr} = 1$ when repetition r of topic-subgroup combination t is active. Activation is linked to the assignments by

$$\begin{aligned} a_{tr} &\geq x_{gtr}, & \forall g, t, r, \\ a_{tr} &\leq \sum_{g=1}^G x_{gtr}, & \forall t, r. \end{aligned} \quad (7)$$

These conditions ensure that a repetition is marked active exactly when at least one self-formed group is assigned to it. This indicator is then used in the repetition and capacity constraints below.

The number of repetitions of each base topic is bounded by

$$\begin{aligned} \sum_{r=1}^{R_t} a_{tr} &\geq r_{\min}, \\ \sum_{r=1}^{R_t} a_{tr} &\leq r_{\max}, \end{aligned} \quad \forall t \in \{1, \dots, T\}. \quad (8)$$

The constraints are written for the first subgroup of each base topic. They require every offered topic to have between $r_{\min}$ and $r_{\max}$ active repetitions.

The number of active repetitions is kept equal across subgroups belonging to the same base topic:

$$\sum_{r=1}^{R_t} a_{tr} = \sum_{r=1}^{R_t} a_{(bT+t)r}, \quad \forall t \in \{1, \dots, T\}, \quad b = 1, \dots, B-1. \quad (9)$$

With topic-subgroup columns ordered as $T_1 S_1, \dots, T_T S_1, T_1 S_2, \dots$, this equality gives every subgroup of a base topic the same number of active repetitions. The limits in (8) therefore propagate to all B subgroups.

Every active topic-subgroup repetition must satisfy its student-capacity bounds:

$$\begin{aligned} \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} &\geq a_{tr} n_{tr}^{\min}, \\ \sum_{i=1}^N \sum_{g=1}^G m_{ig} x_{gtr} &\leq a_{tr} n_{tr}^{\max}, \end{aligned} \quad \forall t \in \{1, \dots, BT\}, r. \quad (10)$$

The membership matrix converts group-level assignments into student counts. Thus, every active repetition must contain at least $n_{tr}^{\min}$ and at most $n_{tr}^{\max}$ students, while inactive repetitions remain empty.

Table 2: Parameter effect guide for the preference-based assignment model (directional, holding other settings fixed).

parameter	if increased	trade-off
n_topics	provides more base topics, each subject to the repetition rules	increases model size and may tighten feasibility
B	adds more subgroups to every topic	permits finer organization but adds balancing requirements
R (R_t)	provides more candidate repetitions	increases model size and only adds flexibility when rmax permits
nmin	requires more students in each active repetition	prevents underfilled groups but tightens feasibility
nmax	permits more students in each repetition	increases flexibility but may permit crowded groups
rmin	requires more repetitions of every topic and corresponding subgroup	tightens feasibility and may reduce achievable preference scores
rmax	permits more repetitions	increases flexibility but may fragment allocations

Multi-role workload allocation

Purpose and inputs

The third model addresses a general multi-role workload allocation problem. Individuals receive discrete work units across several roles while the allocation satisfies task demand, annual workload requirements, and policy constraints. The implementation is motivated by PhD workload allocation, where students receive teaching assistantship, grading, and residual support duties. More generally, role A is preference-sensitive, role B is routine or demand-driven, and role C is allocated according to a priority score such as seniority, experience, or eligibility.

The extractor aligns a student table, preference matrix, and role-demand matrix. The student table supplies cohort year and previous TA and grading workload, which become the priority score q_i and historical workload vectors $t_i^{(1)}$ and $g_i^{(1)}$. The preference matrix supplies P_{ij} , and the demand matrix supplies d_{jr} . Residual-role demand may either be provided or derived from total semester workload and submitted TA and grading demand.

Objective

Let X_{ijr} be the non-negative integer units of role r assigned to individual i for task j . Let Y_1 denote the protected subgroup and Y_0 the individuals included in the role- A fairness calculation. The objective is

$$\min \quad \alpha(L_{\max} - L_{\min}) - \beta \sum_{i,j} P_{ij} X_{ij,A} - \phi \sum_{i,j} q_i X_{ij,C} + \rho \sum_{i \in Y_1} w_i.$$

The first term penalizes the spread in cumulative role- A workload among individuals in Y_0 . The second rewards preference fit for role A , the third rewards assigning role C according to priority q_i , and the final term penalizes violations of the protected subgroup's soft role- A limit. Role B is governed by demand and workload constraints rather than a direct objective reward.

Key constraints

Every task-role demand must be met exactly:

$$\sum_i X_{ijr} = d_{jr}, \quad \forall j, r. \quad (11)$$

Each individual's annual workload is fixed at $2C$:

$$\left(t_i^{(1)} + \sum_j X_{ij,A} \right) + \left(g_i^{(1)} + \sum_j X_{ij,B} \right) + \sum_j X_{ij,C} = 2C, \quad \forall i. \quad (12)$$

The fairness variables bound cumulative role- A workload only for the population included in the spread calculation:

$$L_{\min} \leq t_i^{(1)} + \sum_j X_{ij,A} \leq L_{\max}, \quad \forall i \in Y_0. \quad (13)$$

For protected individuals, w_i permits but penalizes role- A workload above the soft threshold:

$$\sum_j X_{ij,A} \leq t_{\max}^{(Y1)} + w_i, \quad \forall i \in Y_1. \quad (14)$$

Optional role-wide lower and upper bounds are imposed only when supplied:

$$\ell_r \leq \sum_j X_{ijr} \leq u_r, \quad \forall i, r. \quad (15)$$

The variable domains are

$$X_{ijr} \in \mathbb{Z}_{\geq 0}, \quad L_{\min}, L_{\max} \geq 0, \quad w_i \geq 0 \quad (i \in Y_1). \quad (16)$$

In the PhD workload example, roles A , B , and C correspond to teaching assistantship, grading, and residual support. The score q_i is seniority, Y_1 contains Year-1 students, and Y_0 contains the remaining




Figure 1: High-level package architecture.

students. Historical workload allows the model to balance past and current-semester duties together. For a one-semester analysis, users can set `past_ta = 0` and `past_gr = C` before calling `extract_info()`.

Setting an objective weight to zero removes the corresponding policy component from the optimization criterion without removing any associated feasibility constraints.

Table 3: Parameter effect guide for the multi-role workload allocation model (directional, holding other settings fixed).

parameter	if increased	trade-off
alpha	increases the penalty on role-A workload spread within Y_0	can improve fairness at the expense of preference or priority terms
beta	increases the reward for role-A preference scores	can improve preference attainment at the expense of workload spread or priority terms
phi	increases the reward for assigning role C to higher-priority individuals	can improve priority alignment at the expense of the other objective terms
rho	makes protected-group slack more expensive	discourages overload but does not alter feasibility because slack remains available
$t_{\max_y1}$	raises the protected subgroup's soft role-A threshold	weakens protection and can reduce the slack penalty; formal feasibility is unchanged
C	raises every individual's required annual workload from $2C$	requires two additional workload units per person for each unit increase and may cause infeasibility if aggregate demand is incompatible
$ta_{\min}, gr_{\min}, e_{\min}$	raises the required minimum for the corresponding role	tightens feasibility and forces more of that role onto every individual
$ta_{\max}, gr_{\max}, e_{\max}$	raises the permitted maximum for the corresponding role	weakly enlarges the feasible set and may allow greater concentration

3 The grouper package

Figure 1 summarizes the architecture of `grouper` as a domain-specific implementation layer for the allocation models above. The pipeline accepts student information together with allocation attributes such as preferences, demand, capacities, and policy constraints. Inside the package, these inputs pass through data processing, model building, and solving. The final stage returns a raw solver result for diagnostics and reproducibility together with an allocation table for operational use.

Wrapper contract

Each stage in Figure 1 is exposed through a high-level wrapper, with model-specific functions available underneath for advanced use. This contract is intentionally stable across the supported assignment settings. Users can run a concise end-to-end pipeline for routine use, while retaining granular control over input preparation, model construction, solver choice, and post-processing when a local policy requires it.

`extract_info()`

The `extract_info()` function standardizes raw inputs into the list structure used by the model builders. For `assignment = "diversity"`, it accepts the student or group data frame, the self-formed-group column, and either a precomputed dissimilarity matrix or demographic columns from which dissimilarity is computed. For `assignment = "preference"`, it accepts the same group structure together with the group-topic preference matrix. For the historical `assignment = "phd"` option, which invokes the multi-role workload workflow, it accepts the student table, preference matrix, and role-demand matrix. In all cases, the purpose of `extract_info()` is to align user-supplied records with the mathematical objects used in the optimization model.

The following illustrative calls use the historical `assignment = "phd"` argument to show the multi-role workload workflow, with `S`, `P`, `D`, and `J` denoting the student table, preference matrix, demand matrix, and course-code vector.

```
# "diversity": dframe + self_formed_groups + d_mat/demographic_cols
# "preference": dframe + self_formed_groups + pref_mat
# "phd": S, P, and D for multi-role workload allocation
df <- extract_info(
  "phd",
  student_df = S,
  p_mat = P,
  d_mat = D
)
```

`prepare_model()`

The `prepare_model()` function converts standardized inputs and policy parameters into an `ompr` mixed-integer optimization model. For diversity-based assignment, it combines the extracted membership and dissimilarity information with topic, repetition, size, and objective-weight parameters. For preference-based assignment, it builds the topic-slot assignment model from group sizes, preference scores, and slot-capacity parameters. For multi-role workload allocation, it passes the extracted workload inputs and policy weights to the role-allocation model. The returned object is a solver-ready model rather than an allocation table, which lets users inspect or modify the model before solving when needed.

```
# "diversity": add n_topics, R, nmin, nmax, rmin, rmax, w1, w2
# "preference": add n_topics, B, R, nmin, nmax, rmin, rmax
# "phd": alpha, beta, phi, rho, t_max_y1, e_max, and C
m <- prepare_model(
  df,
  assignment = "phd",
  alpha = a,
  beta = b,
  phi = p,
  rho = r
)
```

`solve_assignment()`

The `solve_assignment()` function solves a prepared model and converts the raw solver solution into an allocation output. The function uses `ompr.roi` to call an ROI-backed solver, with `solver = "glpk"`, `"highs"`, or `"gurobi"` selecting the backend when the corresponding plugin or solver is available. It returns a list with `model_result`, the raw output from `ompr::solve_model()` for diagnostics and reproducibility, and `output`, the post-processed assignment table. For the student-group workflows,

the output maps groups back to selected topics or slots; for the multi-role workload workflow, the output is a wide individual-by-task-role allocation table.

```
# solver can be "glpk", "highs", or "gurobi" when available
# group workflows pass dframe, params_list, and group_names
# multi-role passes student_df and course_codes for post-processing
out <- solve_assignment(
  m,
  "phd",
  "glpk",
  student_df = S,
  course_codes = J
)
```

Conceptually, all supported workflows share the same pipeline contract. Institutional data are converted into standardized model inputs and policy choices are translated into constraints and objective terms. Solver results are converted into interpretable allocation outputs. This contract is the main software contribution of grouper relative to lower-level modeling toolkits such as ompr, ROI, and solver interfaces. Those tools provide flexible optimization infrastructure, whereas grouper provides a domain-specific layer that connects higher-education data, institutional policy, model construction, and decision-ready allocation outputs.

4 Multi-role workload allocation in practice

This section evaluates the multi-role workload workflow using institutional planning data from AY2420, AY2510, and AY2520. The notation follows the department's academic-year convention, where AY denotes academic year, the first two digits indicate the academic year, and the final two digits indicate the semester. For example, AY2510 refers to Semester 1 of Academic Year 2025–2026, while AY2520 refers to Semester 2 of the same academic year. Together, the three cases cover different cohort sizes, course offerings, and role-demand profiles, allowing us to assess the workflow across multiple operational settings and compare optimized schedules with the manual schedules used by the department. Student identifiers and course codes are anonymized throughout.

Table 4 summarizes the size and structure of each semester instance. Students and Courses describe the scale of the assignment problem, while TA, GR, and E report the required units for teaching assistantship, grading, and residual support, respectively. In this use case, TA is the preference-sensitive role, GR is primarily demand-driven, and E is assigned according to seniority. The columns Y1–Y4 show the cohort composition by year. AY2510 is the largest instance by both student and course count, while role demand and cohort structure vary across the three semesters.

Table 4: Scale, role demand, and cohort composition of the three multi-role workload datasets.

Semester	Students	Courses	TA	GR	E	Y1	Y2	Y3	Y4
AY2420	43	19	54	80	38	10	13	8	12
AY2510	48	25	50	104	38	15	12	11	10
AY2520	40	19	55	81	24	10	11	13	6

The same wrapper workflow is applied to all three semesters. To avoid repeating nearly identical code, AY2520 is used only to demonstrate the input, model-building, and solving steps and to show one student-level workload distribution.

All three cases are formulated as two-semester workload problems. The current semester is optimized, while previous-semester TA and grading loads are supplied as `past_ta` and `past_gr`. This lets the model balance annual workload rather than treating each semester in isolation.

The workflow begins by standardizing the inputs. In this example, `student_df` records the anonymized student identifier, cohort year, and previous-semester TA/GR workload. The preference matrix P is student by course, and the demand matrix D is course by role. Residual support demand is already included as the third role column, so `e_mode` is set to "none".

```
df_phd <- extract_info(
  assignment = "phd",
  student_df = students_ay2520[, c("student_id", "year", "past_ta", "past_gr")],
  p_mat = P_ay2520,
```

```

d_mat = D_ay2520,
e_mode = "none",
C = 4
)

```

The resulting `df_phd` object is the standardized input contract used by the model-building stage. It contains the aligned preference matrix, role-demand matrix, seniority vector, and past workload vectors; downstream code no longer needs to know the original file layout.

The second step constructs the mixed-integer model. In this example, the policy weights are set to $\alpha = 2$, $\beta = 1$, $\phi = 1$, and $\rho = 10$, encoding the trade-off between yearly TA fairness, TA preference fit, seniority-guided residual support work, and Year-1 TA overload protection. The bound $t_{\max_y1} = 1$ limits Year-1 teaching assistantship load before slack is used, and $e_{\max} = 1$ keeps residual support duties from concentrating too heavily on any single student.

```

model_phd <- prepare_model(
  df_list = df_phd,
  assignment = "phd",
  alpha = 2,
  beta = 1,
  phi = 1,
  rho = 10,
  t_max_y1 = 1,
  e_max = 1,
  C = 4
)

```

At this point, the optimization model is fully specified but not yet solved. The third step selects the solver and requests the operational output mapping. GLPK is used here because it is open-source and supports a reproducible example in a standard R environment.

```

phd_solution <- solve_assignment(
  model = model_phd,
  assignment = "phd",
  solver = "glpk",
  student_df = students_ay2520,
  course_codes = demand_ay2520$course_code,
  name_col = "student_id",
  verbose = FALSE
)

```

The returned object contains `model_result`, which keeps the raw solver result for diagnostics, and output, which stores the allocation in a wide student-by-course-role table. The results are first summarized across all three semesters by comparing the optimized schedules with the manual schedules used by the department. Both are evaluated with the same objective expression. Because the model minimizes a weighted objective with no single natural unit, values should be compared only between schedules evaluated with the same policy weights. Figure 2 plots the absolute objective value for the model and manual schedules so that all bars share a common positive orientation. All six raw values are negative; therefore, a taller bar represents a more negative, and hence lower, raw objective.




Figure 2: Model and manual schedules compared by absolute objective value across three semesters.

The model obtains a lower weighted objective in all three semesters. This should be interpreted as a policy-consistent comparison rather than a general judgment on the department's manual allocation. Both schedules are scored using the same objective function, so each difference reflects the combined treatment of workload fairness, preference fit, seniority steering, and Year-1 protection under the stated policy weights.

Figure 3 then shows how an optimized schedule is allocated at student level. For AY2520, it displays annual workload by student and cohort year. The lower part of each stacked bar is the optimized current-semester allocation, while the upper part is previous-semester workload carried into the annual balance.




Figure 3: AY2520 model workload distribution by anonymized student and cohort year.

The dashed horizontal line marks the four-unit semester capacity. Since each bar combines current and past workload, bars reaching eight units represent a full two-semester workload target. Facing by cohort year makes the Year-1 protection and seniority-based residual-support pattern easier to inspect. Year-1 students receive at most one current-semester TA unit, while residual-support units are assigned only to students in Years 2–4. Among Years 2–4, annual TA workload ranges from one to four units, giving the three-unit spread reported for the baseline solution.

Figure 4 examines the semester with the largest model–manual objective gap, AY2420, by decomposing the gap into weighted objective terms.

The decomposition makes the trade-offs explicit. In AY2420, the preference term accounts for 125 objective units of the 125-unit total gap. The fairness term adds 8 units to the manual objective, while the seniority-residual-support term offsets the gap by 8 units and the Year-1 slack term is unchanged. In this sense, grouper makes allocation rules explicit and searches the feasible assignment space more systematically than an operational manual process.




Figure 4: AY2420 weighted objective-term differences between manual and model schedules.

Hyperparameter effectiveness

A one-at-a-time AY2520 sensitivity check was used to examine whether local changes to the objective weights alter the operational conclusions. In this example, changing alpha has the clearest effect on TA workload fairness because it directly penalizes the annual TA spread. In the local check, the spread decreases from 5 units at $\alpha = 1$ to 3 units at the manuscript value $\alpha = 2$, and remains 3 units at $\alpha = 4$. Local changes to beta, phi, and rho have smaller visible effects because the baseline solution already achieves high TA preference attainment, assigns residual support toward senior students, and uses no Year-1 TA-overload slack. The sensitivity check therefore supports interpreting these weights as local policy levers rather than universal tuning constants.

The reproducible benchmark repeated the AY2520 solve 30 times for each open-source solver. GLPK returned the same objective value, -182 , in all 30 runs with a mean runtime of 0.190 seconds and an observed range from 0.184 to 0.242 seconds. HiGHS also returned objective value -182 in all 30 runs, with a mean runtime of 0.201 seconds and an observed range from 0.192 to 0.248 seconds. These timings indicate that the example is small enough for either open-source backend to solve interactively.

5 Web frontend

grouper also ships with a shiny (Chang et al., 2024) frontend for users who prefer a browser-based workflow over direct R scripting. The unified application is stored under `inst/shiny/grouperWebApp/` and is built with `bslib` (Sievert et al., 2025) and interactive tables from DT (Xie et al., 2024). It is intended for instructors and administrators who need to prepare allocation inputs, set policy parameters, run an optimization model, inspect outputs, and export reusable assignment files without writing the wrapper calls by hand.

The following command launches the application locally.

```
library(shiny)
runApp(system.file("shiny", "grouperWebApp", "app.R"),
       package = "grouper")
```

The same application can also be hosted as an internal website, for example on a Shiny Server or similar institutional platform. This allows educators to make the allocation interface available to course coordinators, administrators, or other staff without requiring each user to install R packages or run optimization code locally. The home tab acts as a model selector. From there, users can open one of the three allocation modules.

The diversity-based assignment (DBA) module is intended for instructors who need to form balanced tutorial, lab, project, or discussion groups from student-level attributes. Users upload the student records, select the self-formed group column, choose demographic and skill variables, tune diversity and size parameters, and export the merged group assignment. The preference-based assignment (PBA) module supports courses where self-formed groups rank topics, supervisors, project streams, or other limited slots; it follows the same pattern of uploading data, setting policy parameters, solving the model, and exporting the resulting assignment.




Figure 5: Home page of the web application.

The multi-role workload tab validates uploaded files, normalizes prior workload, and then runs the same multi-role optimization model described in the workflow section. After a successful run, the app displays solver status, objective diagnostics, a workload-distribution plot, preference attainment, and the wide assignment table produced by `assign_job()`.

The assignment table can be downloaded as an XLSX file for further processing or record-keeping for all three workflows.



Figure 6: Multi-role workload results view.

The frontend is a convenience layer over the same package functions rather than a separate model implementation. Solver availability still depends on installed ROI plugins, and Gurobi-specific time and iteration limits only apply when the Gurobi backend is selected. For deployed use, uploaded files may contain sensitive student information, so the application should be run in an institutionally controlled environment with appropriate access restrictions.

6 Discussion

grouper is designed for higher-education allocation settings where policy rules are explicit, repeated across offerings, and difficult to manage reliably by hand. Its main advantage is that it combines three policy-aware mathematical models with a common workflow over general-purpose optimization tools. Compared with spreadsheet-based or ad hoc procedures, the package makes allocation rules explicit, records the chosen policy parameters, and returns outputs that can be mapped back to teaching and administrative records. Users nevertheless retain control over priorities such as diversity, preference satisfaction, workload balance, and protected-group constraints.

This design also creates practical limitations. Because the package encodes policy choices explicitly, it requires clean and consistently aligned inputs. Student rows, preference matrices, demand matrices, group identifiers, and topic or course labels must refer to the same ordering before the model is built. Richer constraints and objective weights increase flexibility, but they also increase the input burden and the need for local calibration. In practice, the objective weights should be treated as policy parameters rather than universal defaults. Institutions should confirm that the resulting trade-offs

reflect their own allocation policies before using a model operationally.

Performance is governed primarily by the size and structure of the resulting mixed-integer program. The package itself is not internally parallelized. Instead, it delegates optimization to ROI-backed solvers through `ompr.roi`; any parallel or threaded execution depends on the selected solver and its configuration. The diversity-based model grows strongly with the number of student pairs because it uses pairwise dissimilarity terms. The preference-based model grows with the number of self-formed groups, topic-subgroup slots, and repetitions. The multi-role workload model grows with the number of individuals, tasks or courses, and role types. The three-semester examples demonstrate interactive use for moderate departmental instances, but larger deployments should be benchmarked with local data, hardware, and solver settings.

The package uses a functional interface rather than defining an S3, S4, or R6 object hierarchy. Inputs are standardized as lists, matrices, and data frames, models are represented by `ompr` objects, and solved allocations are returned as data frames. This keeps intermediate objects inspectable, although it provides less formal encapsulation than a class-based design. The dependency structure follows the same principle of using established tools for specialized tasks. `ompr` supplies the algebraic modeling layer, while `ompr.roi` and ROI plugins provide access to solvers such as GLPK, HiGHS, and Gurobi. Packages including `cluster`, `dplyr`, and `yaml` support distance computation, data handling, and configuration parsing. This modular design avoids embedding a custom solver in `grouper`, but solver availability, installation, and licensing remain deployment considerations.

7 Conclusion

`grouper` provides three mathematical models and a consistent R workflow for recurring allocation problems in higher education. Diversity-based assignment supports balanced group composition, preference-based assignment allocates self-formed groups to requested destinations, and multi-role workload allocation combines demand, preference, fairness, and priority rules. Across these settings, the package translates institutional policies into explicit model terms exposed through documented R interfaces.

The detailed multi-role application demonstrates this design on a departmental planning task across three semesters. The workflow combines prior workload, course demand, student preferences, seniority-based priority, and Year-1 protection. Under the stated policy weights, the optimized schedules obtain lower objective values than the corresponding manual schedules. This comparison is specific to the encoded policy, while the first two workflows extend the same data-build-solve pattern to diversity balancing and preference-sensitive topic allocation.

Overall, `grouper` fills a practical gap between low-level optimization toolkits and the allocation decisions routinely faced by instructors and administrators. Its scripted and Shiny interfaces support both reproducible analysis and operational use. Future development will focus on extending the supported allocation models, improving diagnostic summaries, and strengthening interfaces for institutional deployment while preserving the common workflow.

Bibliography

- A. A. Anwar and A. S. Bahaj. Student project allocation using integer programming. *IEEE Transactions on Education*, 46(3):359–367, 2003. doi: 10.1109/TE.2003.813519. URL <https://doi.org/10.1109/TE.2003.813519>. [p1, 2]
- W. Chang, J. Cheng, J. Allaire, C. Sievert, B. Schloerke, Y. Xie, J. Allen, J. McPherson, A. Dipert, and B. Borges. *shiny: Web Application Framework for R*, 2024. URL <https://CRAN.R-project.org/package=shiny>. R package version 1.10.0. [p12]
- Free Software Foundation. *GLPK (GNU Linear Programming Kit)*, 2026. URL <https://www.gnu.org/software/glpk/>. [p2]
- Gurobi Optimization, LLC. *Gurobi Optimizer Reference Manual*, 2026. URL <https://www.gurobi.com/documentation/>. [p2]
- HiGHS Developers. *HiGHS: High Performance Software for Linear Optimization*, 2026. URL <https://highs.dev/>. [p2]
- K. Hornik, S. Theussl, F. Schwendinger, M. Schwenk, and G. Stigler. *ROI: R Optimization Infrastructure*, 2024. URL <https://CRAN.R-project.org/package=ROI>. R package version 1.0-1. [p1]

- H. Kim and C. Nolte. *golfr: Group Assignment Tool*, 2025. URL <https://CRAN.R-project.org/package=golfr>. R package version 0.1.0. [p2]
- D. Meyer. Optassign—a web-based tool for assigning students to groups. *Computers & Education*, 53(4): 1104–1119, 2009. doi: 10.1016/j.compedu.2009.04.014. URL <https://doi.org/10.1016/j.compedu.2009.04.014>. [p1, 2]
- J. Ramotsisi, M. Kgomotso, and L. Seboni. An optimization model for the student-to-project supervisor assignment problem: The case of an engineering department. *Journal of Optimization*, 2022(1): 9415210, 2022. doi: 10.1155/2022/9415210. URL <https://doi.org/10.1155/2022/9415210>. [p1]
- D. Schumacher. *ompr: Model and Solve Mixed Integer Linear Programs*, 2025a. URL <https://CRAN.R-project.org/package=ompr>. R package version 1.0.4. [p1]
- D. Schumacher. *ompr.roi: ROI Plugin for ompr*, 2025b. URL <https://CRAN.R-project.org/package=ompr.roi>. R package version 1.0.2. [p1]
- C. Sievert, W. Chang, and J. Cheng. *bslib: Custom Bootstrap Sass Themes for shiny and rmarkdown*, 2025. URL <https://CRAN.R-project.org/package=bslib>. R package version 0.9.0. [p12]
- Y. Xie, J. Cheng, and X. Tan. *DT: A Wrapper of the JavaScript Library DataTables*, 2024. URL <https://CRAN.R-project.org/package=DT>. R package version 0.33. [p12]

Mingyuan Zhang
National University of Singapore
Department of Statistics and Data Science
Singapore
ORCID: 0009-0009-1814-4962
mingyuan.z@u.nus.edu

Kevin Lam
National University of Singapore
Department of Computer Science
Singapore
ORCID: 0000-0002-1535-7059
lamfy@u.nus.edu

Vik Gopal
National University of Singapore
Department of Statistics and Data Science
Singapore
ORCID: 0000-0002-7232-2063
vik.gopal@u.nus.edu